



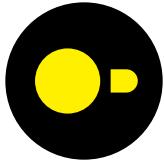




DuckDB

dplyr

Kirill Müller, R/Pharma 2025



DuckDB

Embedded Database

Data: flights X

Clear Column Sorting Convert to Code

dep_time is not missing month > 1 +

Sort by Original Filter

year month day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay	carrier
> # year	111296	2013	2	1	456.00	500	-4.00	652.00	648	4.00 US
> # month	111297	2013	2	1	520.00	525	-5.00	816.00	820	-4.00 UA
> # day	111298	2013	2	1	527.00	530	-3.00	837.00	829	8.00 UA
> # dep_time	111299	2013	2	1	532.00	540	-8.00	1007.00	1017	-10.00 B6
> # sched_dep_time	111300	2013	2	1	540.00	540	0.00	859.00	850	9.00 AA
> # dep_delay	111301	2013	2	1	552.00	600	-8.00	714.00	715	-1.00 EV
> # arr_time	111302	2013	2	1	552.00	600	-8.00	919.00	910	9.00 AA
> # sched_arr_time	111303	2013	2	1	552.00	600	-8.00	655.00	789	-14.00 B6
> # arr_delay	111304	2013	2	1	553.00	600	-7.00	833.00	815	18.00 FL
> # arr_delay	111305	2013	2	1	553.00	600	-7.00	821.00	825	-4.00 MQ
> # arr_delay	111306	2013	2	1	553.00	600	-7.00	659.00	659	0.00 US
> # arr_delay	111307	2013	2	1	554.00	600	-6.00	713.00	716	-3.00 EV
Missing	111308	2013	2	1	554.00	600	-6.00	851.00	984	-13.00 B6
Min	111309	2013	2	1	554.00	601	-7.00	920.00	918	2.00 UA
Median	111310	2013	2	1	555.00	600	-5.00	903.00	906	-3.00 B6
Mean	111311	2013	2	1	556.00	600	-4.00	730.00	745	-15.00 AA
Max	111312	2013	2	1	557.00	600	-3.00	859.00	859	0.00 B6
SD	111313	2013	2	1	558.00	600	-2.00	916.00	912	4.00 B6
> A carrier	111314	2013	2	1	558.00	600	-2.00	738.00	759	-21.00 DL
> # flight	111315	2013	2	1	558.00	605	-7.00	753.00	805	-12.00 MQ
> A tailnum	111316	2013	2	1	559.00	600	-1.00	923.00	925	-2.00 UA
> A origin	111317	2013	2	1	600.00	600	0.00	833.00	837	-4.00 DL
> A dest	111318	2013	2	1	601.00	600	1.00	717.00	730	-13.00 WN
> # air_time	111319	2013	2	1	601.00	608	-7.00	703.00	725	-22.00 UA
> # distance	111320	2013	2	1	601.00	608	-7.00	723.00	755	-32.00 UA
> # hour	111321	2013	2	1	602.00	600	2.00	655.00	658	-3.00 US
> # minute	111322	2013	2	1	603.00	610	-7.00	913.00	915	-2.00 AA
> time_hour	111323	2013	2	1	604.00	610	-6.00	752.00	817	-25.00 DL
> # air_time	111324	2013	2	1	604.00	615	-11.00	747.00	750	-3.00 MQ
> # distance	111325	2013	2	1	605.00	610	-5.00	719.00	735	-16.00 WN
> # hour	111326	2013	2	1	608.00	615	-7.00	837.00	842	-5.00 DL
> # minute	111327	2013	2	1	609.00	610	-1.00	902.00	902	0.00 B6
> time_hour	111328	2013	2	1	610.00	615	-5.00	905.00	855	10.00 9E
> time_hour	111329	2013	2	1	612.00	600	12.00	714.00	703	11.00 US
> time_hour	111330	2013	2	1	615.00	610	5.00	1051.00	1051	0.00 B6

Showing 302,038 rows (89.69% of 336,776 total) 19 columns

DuckDB vs. SQLite

Embedded

Database

Analytical

DuckDB

- . SQLite for analytics
- . Command line interface
- . Python/Rust/WASM/... package
- . R package

Large



Fast



Featureful







DuckDB interfaces

dplyr-like



SQL

DBI



DuckDB

```
SELECT COUNT(*)  
FROM 'personas.parquet';
```

```
--
```

```
--
```

```
--
```

```
--
```

```
--
```

```
--
```

count_star()
int64

100000



```
tbl <- duckdb::tbl_file(  
  path = "personas.parquet")  
tbl |> dplyr::count()  
#> # Source: SQL [?? x 1]  
#> # Database: DuckDB 1.4.1  
#>      n  
#>    <dbl>  
#> 1 100000
```

```
library(DBI)  
con <- dbConnect(duckdb::duckdb())  
dbGetQuery(  
  con,  
  "SELECT COUNT(*)  
   FROM 'personas.parquet'")  
#>   count_star()  
#> 1 1e+05
```

```
SELECT COUNT(*)  
FROM 'personas.parquet';
```

--	count_star()
--	int64
--	100000
--	

```
tbl <- duckdb::tbl_file(  
  path = "personas.parquet")  
tbl |> dplyr::count()  
#> # Source:   SQL [?? x 1]  
#> # Database: DuckDB 1.4.1  
#>      n  
#>    <dbl>  
#> 1 100000
```

```
tbl <- duckdb::read_parquet_duckdb(  
  "personas.parquet")  
tbl |> dplyr::count()  
#> # A duckdb data frame: 1 variable  
#>   n  
#>   <int>  
#> 1 100000
```

```
library(DBI)
con <- dbConnect(duckdb::duckdb())
dbGetQuery(
  con,
  "SELECT COUNT(*)
   FROM 'personas.parquet'")
#> count_star()
#> 1 1e+05
```

```
SELECT COUNT(*)  
FROM "personas.parquet";
```

count_star()
int64
100000

```
library(duckpyr)

personas <- read_parquet_duckdb(
  "s3://my-bucket/.../*.parquet"
)

personas_count <-
  personas |>
  count(sex, edu)

personas_count |>
  explain()
```

ORDER_BY
sex A, education_level A
~14 rows

HASH_GROUP_BY
Groups:
#0, #1
Aggregates:
count_star()
~100,000 rows

READ_PARQUET
Projections:
sex, education_level
~100,000 rows

```
library(duckdb)

personas <- tibble(
  sex = rep(1:2, each = 7),
  edu = rep(1:7, times = 2))

personas_count <-
  personas |>
  count(sex, edu)

personas_count |>
  explain()
```

ORDER_BY
sex A, edu A
14 rows

HASH_GROUP_BY
Groups:
#0, #1
Aggregates:
count_star()
14 rows

R_DATAFRAME_SCAN
Projections:
sex, edu
14 rows

```
library(tidyverse)

G <- 200
N <- 5000

data <-
  crossing(g1 = 1:G, g2 = 1:G, id = 1:N) |>
  mutate(x = rnorm(n()))

data |>

  summarize(.by = c(g1, g2), m = mean(x)) |>

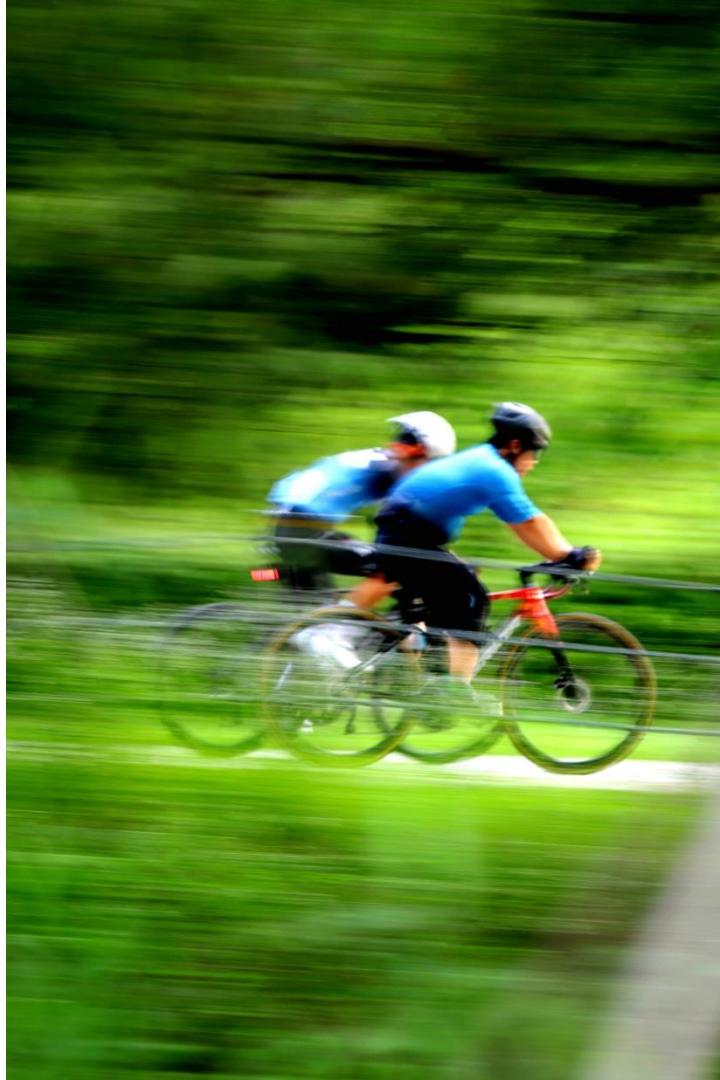
  ...

#>    user  system elapsed
#>    2.397   1.720   5.172
```



Existing pipelines

- All in
- Best effort
- Targeted



```
library(tidyverse)
library(duckplyr)
#> ✓ Overwriting dplyr methods with duckplyr methods.
#> i Turn off with `duckplyr::methods_restore()`.
```

```
G <- 200
N <- 5000
```

```
data <-
  crossing(g1 = 1:G, g2 = 1:G, id = 1:N) |>
  mutate(x = rnorm(n()))
```

```
data |>

  summarize(.by = c(g1, g2), m = mean(x)) |>
```

```
...
#>    user  system elapsed
#>    4.073   0.013   0.544
```



```
library(tidyverse)
library(duckplyr)
#> ✓ Overwriting dplyr methods with duckplyr methods.
#> i Turn off with `duckplyr::methods_restore()`.
methods_restore()
#> i Restoring dplyr methods.
```

```
data <-
  crossing(g1 = 1:G, g2 = 1:G, id = 1:N) |>
  mutate(x = rnorm(n()))

data |>
  as_duckdb_tibble(prudence = "stingy") |>
  summarize(.by = c(g1, g2), m = mean(x)) |>
  collect() |>
  ...
#>    user  system elapsed
#>    4.142   0.017   0.659
```



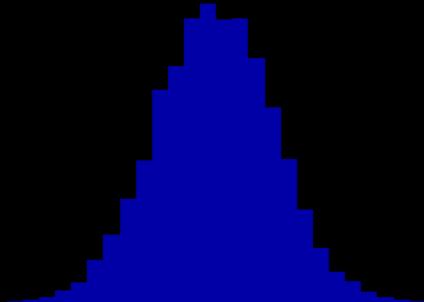
duckpylr

- Output: Proper data frames
- Input: files, data frames, ...
- Fallback, translation
- Full DuckDB power

```
library(tidyverse)  
  
idx <- duckdb::duckdb_tibble(id = 1:1e4)  
  
idx |>  
  mutate(x = dd::random()) |>  
  ggplot(aes(x)) +  
  geom_histogram()
```



```
library(tidyverse)  
  
idx <- duckplyr::duckdb_tibble(id = 1:1e4)  
  
duckplyr::db_exec("LOAD stochastic")  
  
idx |>  
  mutate(x = dd::dist_normal_sample(0, 1)) |>  
  ggplot(aes(x)) +  
  geom_histogram()
```



```
library(duckplyr)

db_exec("INSTALL read_stat")
db_exec("LOAD read_stat")

survey <- read_file_duckdb(
  "survey.sas7bdat",
  "read_stat"
)

survey |>
  count(SEX)
## # A duckplyr data frame
## ...
```



```
install.packages("duckplyr")
library(duckplyr)
methods_restore()
db_exec("INSTALL read_stat")
```



tinyurl.com/duckplyr-25-rp



cynkra